

Programming Assignment 2 - Titanic Dataset

June 25, 2022

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

1 Importing Data

```
[2]: titanic_train = pd.read_csv('train.csv')
titanic_test = pd.read_csv('test.csv')
titanic_test.head()
```

```
[2]:
```

	PassengerId	Pclass	Name	Sex	\
0	892	3	Kelly, Mr. James	male	
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	
2	894	2	Myles, Mr. Thomas Francis	male	
3	895	3	Wirz, Mr. Albert	male	
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	34.5	0	0	330911	7.8292	NaN	Q
1	47.0	1	0	363272	7.0000	NaN	S
2	62.0	0	0	240276	9.6875	NaN	Q
3	27.0	0	0	315154	8.6625	NaN	S
4	22.0	1	1	3101298	12.2875	NaN	S

```
[3]: titanic_train.PassengerId.size
```

```
[3]: 891
```

```
[4]: titanic_train.head()
```

```
[4]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	

```
4          5          0          3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[5]: titanic_train.dtypes
```

```
[5]: PassengerId    int64
Survived          int64
Pclass            int64
Name              object
Sex               object
Age              float64
SibSp             int64
Parch             int64
Ticket            object
Fare              float64
Cabin             object
Embarked          object
dtype: object
```

Check quality of data

```
[6]: titanic_train.isnull().sum()
```

```
[6]: PassengerId    0
Survived          0
Pclass            0
Name              0
Sex               0
Age              177
SibSp             0
Parch             0
Ticket            0
Fare              0
Cabin             687
```

```
Embarked      2
dtype: int64
```

```
[7]: titanic_test.PassengerId.size,titanic_test.isnull().sum()
```

```
[7]: (418,
      PassengerId      0
      Pclass           0
      Name             0
      Sex              0
      Age              86
      SibSp            0
      Parch            0
      Ticket           0
      Fare             1
      Cabin            327
      Embarked         0
      dtype: int64)
```

2 Data Cleaning

Both the training set and test set contain null values in the Age and Cabin columns. The training set contains missing values for the Embarked and for the test set the Fare. Cabin will be dropped from as the majority of the passengers do not have their cabin listed. For the Age and Fare the median will be used. For Embarked the mode will be used. Other columns to be dropped are passengerID and Ticket.

```
[8]: data_raw = [titanic_train,titanic_test]

for df in data_raw:

    df['Age'].fillna(df['Age'].median(), inplace = True)

    df['Embarked'].fillna(df['Embarked'].mode()[0], inplace = True)

    df['Fare'].fillna(df['Fare'].median(), inplace = True)

drop_column = ['PassengerId','Cabin', 'Ticket']
titanic_train.drop(drop_column, axis=1, inplace = True)
```

Sex and Embarked are converted to integers to be used in the classifier.

```
[9]: sex_conv = {'male':0, 'female':1}
titanic_train['Sex'] = titanic_train['Sex'].map(sex_conv)
titanic_test['Sex'] = titanic_test['Sex'].map(sex_conv)

titanic_train.head()
```

```
[9]:
```

	Survived	Pclass		Name	Sex	\
0	0	3		Braund, Mr. Owen Harris	0	
1	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...		1	
2	1	3		Heikkinen, Miss. Laina	1	
3	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)		1	
4	0	3		Allen, Mr. William Henry	0	

	Age	SibSp	Parch	Fare	Embarked
0	22.0	1	0	7.2500	S
1	38.0	1	0	71.2833	C
2	26.0	0	0	7.9250	S
3	35.0	1	0	53.1000	S
4	35.0	0	0	8.0500	S

```
[10]: emb_conv = {'S':0, 'C':1, 'Q':2}
titanic_train['Embarked'] = titanic_train['Embarked'].map(emb_conv)
titanic_test['Embarked'] = titanic_test['Embarked'].map(emb_conv)

titanic_train.head()
```

```
[10]:
```

	Survived	Pclass		Name	Sex	\
0	0	3		Braund, Mr. Owen Harris	0	
1	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...		1	
2	1	3		Heikkinen, Miss. Laina	1	
3	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)		1	
4	0	3		Allen, Mr. William Henry	0	

	Age	SibSp	Parch	Fare	Embarked
0	22.0	1	0	7.2500	0
1	38.0	1	0	71.2833	1
2	26.0	0	0	7.9250	0
3	35.0	1	0	53.1000	0
4	35.0	0	0	8.0500	0

```
[11]: titanic_train.dtypes
```


```
[11]: Survived      int64
Pclass          int64
Name            object
Sex             int64
Age            float64
SibSp           int64
Parch           int64
Fare            float64
Embarked        int64
dtype: object
```

3 Data Exploration

My first thought was that those who are in a higher class cabin had a better chance of survival. Plotted the count of survivors (=1) and non-survivors(=0) by class.

```
[12]: sns.countplot(x='Pclass',data = titanic_train, hue = 'Survived')
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x192a731c400>
```



Looking at actual survival rates.


```
[13]: titanic_train[['Pclass', 'Survived']].groupby(['Pclass']).mean().  
      ↪sort_values(by='Survived')
```

```
[13]:      Survived  
Pclass  
3      0.242363  
2      0.472826  
1      0.629630
```

Being in a higher class increased your chances of survival. The same exercise is done for gender, where male = 0 and female = 1

```
[14]: sns.countplot(x='Sex',data=titanic_train, hue='Survived')
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x192a72da4f0>
```




```
[15]: titanic_train[['Sex', 'Survived']].groupby(['Sex']).mean().  
      ↪sort_values(by='Survived')
```

```
[15]:      Survived  
Sex  
0    0.188908  
1    0.742038
```

Being female greatly increased the chances of survival. Number of siblings was looked at next

```
[16]: sns.countplot(x='SibSp',data=titanic_train, hue='Survived')
```



```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x192a7af2400>
```



No strong conclusions can be made with this data. Age was then looked at. First on two different histograms.

```
[17]: titanic_train_survive = titanic_train[titanic_train['Survived'] == 1]
      titanic_train_died = titanic_train[titanic_train['Survived'] == 0]


      titanic_train_survive.Age.hist(bins = 10)
      plt.title("Count of Survival By Age")
      plt.show()
      titanic_train_died.Age.hist(bins = 10)
      plt.title("Count of Death By Age")
      plt.show()
```



And then overlayed to compare distribution



```
[18]: fig,ax = plt.subplots()
sns.histplot(titanic_train[titanic_train['Survived']==0]['Age'],color = 'r',label = 'Died',ax=ax,kde = True)
sns.histplot(titanic_train[titanic_train['Survived']==1]['Age'],color = 'b',label = 'Survived',ax=ax,kde = True)
plt.legend()
```

[18]: <matplotlib.legend.Legend at 0x192a7ca3370>



Being younger increased the change of survival. A pairplot was created to see if there are any other conclusions to be made.


```
[19]: g = sns.pairplot(titanic_train,diag_kind="hist",hue = 'Survived',dropna = True)
```



The pair plot shows that spending more on the fare also increase survival chances. A heatmap was created to see if there are any correlated variables that can be removed.

```
[20]: titanic_train_corr = titanic_train.corr()
sns.heatmap(titanic_train_corr,annot=True)
#print(titanic_train_corr)
#sns.heatmap(titanic_train_corr, fmt='g', cmap = 'Spectral').set(title = '
↳ 'Correlation Matrix')
#plt.show()
```

```
[20]: <matplotlib.axes._subplots.AxesSubplot at 0x192aae17430>
```



Heatmap confirms the high correlation between Sex and Survived. Other correlations are not as strong. PCA was run to determine the variables that have more weight in determining survival.

```
[21]: #pca analysis
from pca import pca


features = ['Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'Sex', 'Embarked']
X = titanic_train.loc[:, features].values
y = titanic_train.loc[:, 'Survived'].values
```

```
[22]: model = pca(n_components=6, normalize = True)
results = model.fit_transform(X, col_labels = features, row_labels = y)

# Plot explained variance
fig, ax = model.plot()
```

```
[pca] >Normalizing input data per feature (zero mean and unit variance)..
[pca] >The PCA reduction is performed on the [7] columns of the input dataframe.
[pca] >Fit using PCA.
[pca] >Compute loadings and PCs.
[pca] >Compute explained variance.
[pca] >Outlier detection using Hotelling T2 test with alpha=[0.05] and
n_components=[6]
```

```
[pca] >Outlier detection using SPE/DmodX with n_std=[2]
```



<Figure size 432x288 with 0 Axes>

```
[23]: print(model.results['topfeat'])
```


	PC	feature	loading	type
0	PC1	Fare	0.571703	best
1	PC2	Age	0.554338	best
2	PC3	Embarked	0.851590	best
3	PC4	Sex	-0.749587	best
4	PC5	Age	0.787430	best
5	PC6	Parch	0.678252	best
6	PC2	Pclass	-0.532689	weak
7	PC6	SibSp	-0.658818	weak

In decending order: Fare, Age, Embarked, and Sex were the top four variables that explained the sample variance. Those four variables are the starting point for the algorithms.

Next is to remove any outliers in the data. Age and Fare are looked at for outliers as the other variables are categorical. An outlier here is defined as a value outside of the 5-95 percentile.


```
[24]: sns.boxplot(x=titanic_train['Age'],whis=[5,95])
```

[24]: <matplotlib.axes._subplots.AxesSubplot at 0x192a72ebe50>



```
[25]: sns.boxplot(x=titanic_train['Fare'],whis=[5,95])
```

[25]: <matplotlib.axes._subplots.AxesSubplot at 0x192ad2fd640>



```
[26]: titanic_train[["Age", "Fare"]].describe(percentiles = [.05,.25,.5,.75,.95])
```

```
[26]:
```

	Age	Fare
count	891.000000	891.000000
mean	29.361582	32.204208
std	13.019697	49.693429
min	0.420000	0.000000
5%	6.000000	7.225000
25%	22.000000	7.910400
50%	28.000000	14.454200
75%	35.000000	31.000000
95%	54.000000	112.079150
max	80.000000	512.329200

```
[27]: titanic_train.drop(titanic_train[titanic_train["Age"] >\
                                (np.percentile(titanic_train["Age"], 95))].
    ↪index, inplace=True)
titanic_train.drop(titanic_train[titanic_train["Age"] <\
                                (np.percentile(titanic_train["Age"], 5))].
    ↪index, inplace=True)

titanic_train.drop(titanic_train[titanic_train["Fare"] >\
                                (np.percentile(titanic_train["Fare"], 95))].
    ↪index, inplace=True)
titanic_train.drop(titanic_train[titanic_train["Fare"] <\
                                (np.percentile(titanic_train["Fare"], 5))].
    ↪index, inplace=True)
```

4 Classification Methods

Splitting the training set for training and validating. A 80/20 split was used.

```
[28]: from sklearn.model_selection import train_test_split

dataCM1 = titanic_train

X = dataCM1.drop('Survived',axis=1)
y = dataCM1['Survived']

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=.2,
    ↪random_state=15)
```

```
[29]: X_train.shape, X_val.shape
```

```
[29]: ((586, 8), (147, 8))
```

```
[30]: X_val.head()
```

```
[30]:
```

	Pclass	Name	Sex	Age	SibSp	Parch	Fare	\
369	1	Aubart, Mme. Leontine Pauline	1	24.0	0	0	69.3000	
502	3	O'Sullivan, Miss. Bridget Mary	1	28.0	0	0	7.6292	
140	3	Boulos, Mrs. Joseph (Sultana)	1	28.0	0	2	15.2458	
384	3	Plotcharsky, Mr. Vasil	0	28.0	0	0	7.8958	
686	3	Panula, Mr. Jaako Arnold	0	14.0	4	1	39.6875	

	Embarked
369	1
502	2
140	1
384	0
686	0

Classification Method 1 - SVC

For the first attempt I am using Support Vector Classification. For the kernel I am using a 5th order polynomial. The important values kept at default are unlimited iterations (`max_iter = -1`) and the stopping criterion (`tol = 1e-3`).

```
[31]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
modelSVC = SVC(kernel = 'poly', degree = 5)

#Determined Important Features
features = ['Fare', 'Age', 'Embarked', 'Sex']

X_trainSVC = X_train.loc[:, features]

modelSVC.fit(X_trainSVC, y_train)

print('Training accuracy = ', accuracy_score(y_train, modelSVC.
↪predict(X_trainSVC)))
#print('Training accuracy = ', modelSVC.score(X_trainSVC, y_train)*100)

X_valSVC = X_val.loc[:, features]
print('Validation accuracy = ', accuracy_score(y_val, modelSVC.
↪predict(X_valSVC)))
```

```
Training accuracy = 0.71160409556314
Validation accuracy = 0.6666666666666666
```

As a second attempt I changed the kernel to Radial Basis Function which performed slightly better.

```
[32]: modelSVC2 = SVC(kernel = 'rbf')

features = ['Fare', 'Age', 'Embarked', 'Sex']

X_trainSVC2 = X_train.loc[:, features]

modelSVC2.fit(X_trainSVC2, y_train)

print('Training accuracy = ', accuracy_score(y_train, modelSVC2.
↪predict(X_trainSVC2)))

X_valSVC2 = X_val.loc[:, features]
print('Validation accuracy = ', accuracy_score(y_val, modelSVC2.
↪predict(X_valSVC2)))
```

```
Training accuracy =  0.6843003412969283
Validation accuracy =  0.6462585034013606
```

Classification Model 2 - Decision Tree

My second attempt used a decision tree classifier. The quality of a split was measured with entropy. The min samples to require a split was increased from the default of 2 to 20. After testing a few values, increasing the amount of samples to initiate a split improved the validation accuracy. It is possible for this dataset that with the default 2 there is excess overfitting to the training set. A greedy approach was used for splitting on the default value splitter="best".

```
[33]: from sklearn.tree import DecisionTreeClassifier
from sklearn import tree

modelDTC = DecisionTreeClassifier(criterion = 'entropy', min_samples_split = 20,
↪splitter = "best")

features = ['Fare', 'Age', 'Embarked', 'Sex']

X_trainDTC = X_train.loc[:, features]

modelDTC.fit(X_trainDTC, y_train)

print('Training accuracy = ', accuracy_score(y_train, modelDTC.
↪predict(X_trainDTC)))
#print('Training accuracy = ', modelDTC.score(X_trainDTC, y_train)*100)

X_valDTC = X_val.loc[:, features]
print('Validation accuracy = ', accuracy_score(y_val, modelDTC.
↪predict(X_valDTC)))
```

```
Training accuracy =  0.8720136518771331
Validation accuracy =  0.7619047619047619
```


Classification Model 3 - Boosting

The third attempt used Gradient Boosting Classifier. The number of boosting stages was increased from 100 to 200. The min split and leaf values were kept as default due to GBCs resilience to over-fitting. The tolerance for stopping was kept at default $\text{tol}=1\text{e-}4$.

```
[34]: from sklearn.ensemble import GradientBoostingClassifier

modelGBC = GradientBoostingClassifier(max_depth = 3, n_estimators = 200)

features = ['Fare', 'Age', 'Embarked', 'Sex']

X_trainGBC = X_train.loc[:, features]

modelGBC.fit(X_trainGBC, y_train)

print('Training accuracy = ', accuracy_score(y_train, modelGBC.
      ↪predict(X_trainGBC)))

X_valGBC = X_val.loc[:, features]
print('Validation accuracy = ', accuracy_score(y_val, modelGBC.
      ↪predict(X_valGBC)))
```

```
Training accuracy = 0.9351535836177475
Validation accuracy = 0.7891156462585034
```

Test Prediction csv creation

```
[35]: titanic_test.head()
```

```
[35]:
```

	PassengerId	Pclass	Name	Sex	\
0	892	3	Kelly, Mr. James	0	
1	893	3	Wilkes, Mrs. James (Ellen Needs)	1	
2	894	2	Myles, Mr. Thomas Francis	0	
3	895	3	Wirz, Mr. Albert	0	
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	1	

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	34.5	0	0	330911	7.8292	NaN	2
1	47.0	1	0	363272	7.0000	NaN	0
2	62.0	0	0	240276	9.6875	NaN	2
3	27.0	0	0	315154	8.6625	NaN	0
4	22.0	1	1	3101298	12.2875	NaN	0

```
[36]: titanic_test.isnull().sum()
```

```
[36]: PassengerId    0
      Pclass        0
      Name          0
```

```
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin         327
Embarked       0
dtype: int64
```

```
[37]: #Prepare Test Data
features = ['Fare', 'Age', 'Embarked', 'Sex']
X_test = titanic_test.loc[:, features].copy()

#Classification Model 1 - SVC
submitSVC = pd.DataFrame(titanic_test['PassengerId'])
submitSVC['Survived'] = modelSVC2.predict(X_test)
submitSVC.to_csv('submitSVC.csv', index = False)

#Classification Model 1 - Decision Tree
submitDTC = pd.DataFrame(titanic_test['PassengerId'])
submitDTC['Survived'] = modelDTC.predict(X_test)
submitDTC.to_csv('submitDTC.csv', index = False)

#Classification Model 1 - Gradient Boosting
submitGBC = pd.DataFrame(titanic_test['PassengerId'])
submitGBC['Survived'] = modelGBC.predict(X_test)
submitGBC.to_csv('submitGBC.csv', index = False)
```

5 Results

```
[38]: from IPython.display import Image
      Image("Titanic_Predictions.png")
```

```
[38]:
```

8 submissions for MichaelB1234		Sort by Select...
All Successful Selected		
Submission and Description		Public Score
submitGBC.csv just now by MichaelB1234 Classification Model 3 - Gradient Boosting		0.77033
submitDTC.csv a few seconds ago by MichaelB1234 Classification Model 2 - Decision Tree		0.72727
submitSVC.csv a minute ago by MichaelB1234 Classification Model 1 - SVC		0.64593

```
[39]: from IPython.display import Image
Image("Titanic_Leaderboard.png")
```

[39]:

9124	zhangxize		0.77033	7	5h
9125	Vitalii8		0.77033	2	4h
9126	Kent Hui		0.77033	4	1h
9127	Masaru Umekawa		0.77033	1	25m
9128	MichaelB1234		0.77033	8	1s
Your Best Entry! Your most recent submission scored 0.77033, which is the same as your previous score. Keep trying!					
9129	Matthew Conger		0.76794	2	2mo
9130	Jason Lam #2		0.76794	4	2mo
9131	Grant Jensen		0.76794	1	2mo
9132	Yash Dubey		0.76794	3	2mo
9133	gokulraj3		0.76794	2	2mo

6 Conclusion

At the time of writing this, my 3rd classification model using Gradient Boosting had an accuracy of 77.03% and placed 9128th on the leader board. Some steps that can be taken to improve accuracy is to change the split of training and validation, modify/optimize all hyperparameters, changing the input variables, and exploring other classification techniques. The way that the NaN for Age are filled can also be changed to use the mean or split the NaN over multiple values using a correlation to another value (classifying the age). The Sex category can be extended to have a 3rd option of ``Child'' as the chance of survival for a child was higher regardless of gender. Looking at the leaderboard there is some success using regressions which could be explored as well.

Other Models Tested

```
[40]: from sklearn.neighbors import KNeighborsClassifier

modelKNN = KNeighborsClassifier(n_neighbors = 3, weights = 'distance')

features = ['Fare', 'Age', 'Embarked', 'Sex']

X_trainKNN = X_train.loc[:, features]

modelKNN.fit(X_trainKNN, y_train)

print('Training accuracy = ', accuracy_score(y_train, modelKNN.
      ↳predict(X_trainKNN)))

X_valKNN = X_val.loc[:, features]
print('Validation accuracy = ', accuracy_score(y_val, modelKNN.
      ↳predict(X_valKNN)))
```

```
Training accuracy = 0.9744027303754266
Validation accuracy = 0.7346938775510204
```

```
[41]: from sklearn.ensemble import RandomForestClassifier

modelRFC = RandomForestClassifier(n_estimators=100, max_depth = 6, criterion = 'gini', n_jobs = -1)

features = ['Fare', 'Age', 'Embarked', 'Sex', 'Parch']

X_trainRFC = X_train.loc[:, features]

modelRFC.fit(X_trainRFC, y_train)

print('Training accuracy = ', accuracy_score(y_train, modelRFC.
      ↳predict(X_trainRFC)))
```

```
X_valRFC = X_val.loc[:, features]
print('Validation accuracy = ', accuracy_score(y_val, modelRFC.
↪predict(X_valRFC)))
```

```
Training accuracy = 0.8737201365187713
Validation accuracy = 0.7278911564625851
```

```
[ ]:
```